


Pandas

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



pandas: a library for data manipulation

► Documentation:

<http://pandas.pydata.org/pandas-docs/stable/>

pandas 0.24.1 documentation »

next | modules | index

Table Of Contents

- What's New in 0.24.1
- Installation
- Getting started
- User Guide
- pandas Ecosystem
- API Reference
- Development
- Release Notes

Search

Enter search terms or a module, class or function name.

pandas: powerful Python data analysis toolkit

Date: Feb 03, 2019 Version: 0.24.1

Download documentation: [PDF Version](#) | [Zipped HTML](#)

Useful links: [Binary Installers](#) | [Source Repository](#) | [Issues & Ideas](#) | [Q&A Support](#) | [Mailing List](#)

pandas is an open source, BSD-licensed library providing high-performance, easy-to-use data structures and data analysis tools for the [Python](#) programming language.

See the [Package overview](#) for more detail about what's in the library.

- What's New in 0.24.1
 - API Changes
 - Fixed Regressions
 - Bug Fixes
 - Contributors
- Installation
 - Plan for dropping Python 2.7
 - Python version support
 - Installing pandas
 - Running the test suite
 - Dependencies
- Getting started
 - Package overview
 - 10 Minutes to pandas
 - Essential Basic Functionality
 - Intro to Data Structures

Reading in any textfile

```
1 import pandas as pd
2 FNAME = "http://www.stat.ucla.edu/projects/datasets/twins.dat"
3
4 df = pd.read_csv(FNAME, sep=',')
```

Inspecting the DataFrame

```
1 df.shape    # Dimensions
2 df.head()   # First 5 lines
3 df.tail()   # Last 5 lines
4 df.columns  # List of variables as list
5 df.describe() # Summary statistics
```

Inspecting the DataFrame

```
1 df.shape    # Dimensions
2 df.head()   # First 5 lines
3 df.tail()   # Last 5 lines
4 df.columns  # List of variables as list
5 df.describe() # Summary statistics
```

1. How many observations do you have?
2. How many variables do you have?
3. Which variables are numeric?
4. What is the mean of variable "DEDUC1"?

Understanding dtypes

```
1 df.info()
```

Understanding dtypes

```
1 df.info()
```

Pandas	Python	Purpose
object	unicode	Text
int64	int	Integers
float64	float	Floating numbers
bool	bool	True and False values
datetime64		Date and time values
timedelta[ns]		Differences between two datetimes
category		Finite list of text values

Slicing the DataFrame

```
1 df["DEDUC"]    # Column by column name
2 df.loc[0]      # Row by index name
3 df.iloc[0]     # Row by row number
4 df[["AGE", "HRWAGEH"]] # Columns by list of names
5 df.loc["AGE":"HRWAGEH"] # Column range by column names (will be empty)
6 df.iloc[:,5:7] # Column range by column indices
7 df.iloc[1, 4]  # Value by row and column index
8 df.loc[18, "WHITEH"] # Value by row and column name
9 df["WHITEH"].iloc[18] # Same as above
```

Slicing the DataFrame

```
1 df["DEDUC"]    # Column by column name
2 df.loc[0]      # Row by index name
3 df.iloc[0]     # Row by row number
4 df[["AGE", "HRWAGEH"]] # Columns by list of names
5 df.loc["AGE":"HRWAGEH"] # Column range by column names (will be empty)
6 df.iloc[:,5:7] # Column range by column indices
7 df.iloc[1, 4]  # Value by row and column index
8 df.loc[18, "WHITEH"] # Value by row and column name
9 df["WHITEH"].iloc[18] # Same as above
```

1. What is the 6th entry of the 5th column?
2. What is the 5th entry of column "DTEN"?
3. What is the last entry of column "HRWAGEL"?

Boolean indexing

```
1 df[df["AGE"] > 20]
2 df[(df["AGE"] > 20) & (df["WHITE"] == 1)]
3 df[~df["AGE"] > 20]
4 df[df["AGE"].isin((20, 21, 22, 23))]
```

Boolean indexing

```
1 df[df["AGE"] > 20]
2 df[(df["AGE"] > 20) & (df["WHITE"] == 1)]
3 df[~df["AGE"] > 20]
4 df[df["AGE"].isin((20, 21, 22, 23))]
```

1. How many observations have "WHITE" equal to 0?
2. How many observations have "WHITE" equal to 1 and "DEDUC1" unequal to 0?
3. What is the mean age of married twins whose L-parent is a non-white male? (Use "DMARRIED", "WHITE" and "MALE")
4. In how many rows do the values for "WHITE" and "WHITE" differ?

Aggregate data

```
1 df["WHITE"].value_counts()
2 pd.crosstab(df["WHITE"], df["WHITE"])
3 df[["DEDUC2", "EDUCL"]].corr()
```

Aggregate data

```
1 df["WHITEH"].value_counts()
2 pd.crosstab(df["WHITEH"], df["WHITEH"])
3 df[["DEDUC2", "EDUCL"]].corr()
```

1. What is the most common value in "EDUCL"?
2. What is the most common combination of "MALEH" and "MALEL"?
3. What is the Spearman correlation between "EDUCH" and "EDUCL"? What is the Pearson correlation? (Check documentation!)

Changing dtypes

```
1 df["DMARRIED"] = df["DMARRIED"].astype(bool)
2 df["WHITEH"] = df["WHITEH"].astype("category")
3 df["HRWAGEH"] = pd.to_numeric(df["HRWAGEH"], errors="coerce")
```

Manipulation

```
1 df = df.sort_values(by='HRWAGEH')  # Sort by column
2 df = df[sorted(df.columns)]        # Re-order columns alphabetically
3 df['new'] = 9  # Add new column
4 df['AGETR'] = df['AGE']**3
5 df['combined'] = df['MALEH'] + df['EDUCH']
6 df["HRWAGEH_new"] = df["HRWAGEH"].fillna(0)  # Fill missings with 0
7 df = df.dropna(subset=["HRWAGEH"])  # Drop rows missing in "HRWAGEH"
```

Grouping

```
1 grouped = df.groupby(['MALEH'])
2 print(grouped['AGE'].mean())
3 print(grouped['EDUCH'].agg(['mean', 'sum']))
4 print(grouped[['EDUCH', 'AGE']].agg(['mean', 'std']))
```

Grouping

```
1 grouped = df.groupby(['MALEH'])
2 print(grouped['AGE'].mean())
3 print(grouped['EDUCH'].agg(['mean', 'sum']))
4 print(grouped[['EDUCH', 'AGE']].agg(['mean', 'std']))
```

→ Full list at [http:](http://pandas.pydata.org/pandas-docs/stable/getting_started/basics.html#descriptive-statistics)

[//pandas.pydata.org/pandas-docs/stable/getting_started/basics.html#descriptive-statistics](http://pandas.pydata.org/pandas-docs/stable/getting_started/basics.html#descriptive-statistics)

- ▶ What is the "AGE" variance for "MALEL" == 0 individuals?
- ▶ What are the second and the third quartile of "EUDCL" for "MALEL" == 0 individuals?
- ▶ What is the average age for individuals with "MALEL" == 0 and "MALEH" = 0?

Creating DataFrames from other objects

Creating Pandas DataFrames from Python Lists and Dictionaries

Row Oriented

Dictionary

```
sales = [{'account': 'Jones LLC', 'Jan': 150, 'Feb': 200, 'Mar': 140},  
        {'account': 'Alpha Co', 'Jan': 200, 'Feb': 210, 'Mar': 215},  
        {'account': 'Blue Inc', 'Jan': 50, 'Feb': 90, 'Mar': 95 }]  
df = pd.DataFrame(sales)
```

default

	account	Jan	Feb	Mar
0	Jones LLC	150	200	140
1	Alpha Co	200	210	215
2	Blue Inc	50	90	95

List

```
sales = [('Jones LLC', 150, 200, 50),  
        ('Alpha Co', 200, 210, 90),  
        ('Blue Inc', 140, 215, 95)]  
labels = ['account', 'Jan', 'Feb', 'Mar']  
df = pd.DataFrame.from_records(sales, columns=labels)
```

from_records

Column Oriented

```
sales = {'account': ['Jones LLC', 'Alpha Co', 'Blue Inc'],  
        'Jan': [150, 200, 50],  
        'Feb': [200, 210, 90],  
        'Mar': [140, 215, 95]}  
df = pd.DataFrame.from_dict(sales)
```

from_dict

```
sales = [{'account', ['Jones LLC', 'Alpha Co', 'Blue Inc']},  
        {'Jan', [150, 200, 50]},  
        {'Feb', [200, 210, 90]},  
        {'Mar', [140, 215, 95]}]  
df = pd.DataFrame.from_items(sales)
```

from_items

When using a dictionary, column order is not preserved.
Explicitly order them:

```
df = df[['account', 'Jan', 'Feb', 'Mar']]
```

Creating DataFrames from other objects, cont.

```
1 d = {'employee': ['Hannes', 'Fabiana', 'George', 'Olga'],
2      'group': ['Accounting', 'Engineering', 'Engineering', 'HR']}
3 df1 = pd.DataFrame.from_dict(d)
4 t = [('Hannes', 2004), ('Fabiana', 2008), ('George', 2012), ('Olga', 2014)]
5 df2 = pd.DataFrame.from_records(t, columns=["employee", "hire_date"])
```

Appending, Concatening and Merging

```
1 df3 = df1.append(df2)
2 df4 = pd.concat([df1, df2])
3
4 df5 = pd.concat([df1, df2], axis=1)
5 df6 = df1.merge(df2, left_on="employee", right_on="employee")
```

Appending, Concatening and Merging

```
1 df3 = df1.append(df2)
2 df4 = pd.concat([df1, df2])
3
4 df5 = pd.concat([df1, df2], axis=1)
5 df6 = df1.merge(df2, left_on="employee", right_on="employee")
```

- ▶ How do objects df3 and df4 differ?
- ▶ How do objects df5 and df6 differ?

What is pivot?

Pivot

df

	foo	bar	baz	zoo
0	one	A	1	x
1	one	B	2	y
2	one	C	3	z
3	two	A	4	q
4	two	B	5	w
5	two	C	6	t



```
df.pivot(index='foo',  
          columns='bar',  
          values='baz')
```

bar	A	B	C
foo			
one	1	2	3
two	4	5	6

from: "Reshaping and Pivot Tables"

Pivoting and melting

```
1 pivoted = df6.pivot(index='employee', columns='group', values='hire_date')
2 reverse = (pivoted.reset_index()
3             .melt(id_vars="employee", value_name="hire_date"))
```

Pivoting and melting

```
1 pivoted = df6.pivot(index='employee', columns='group', values='hire_date')
2 reverse = (pivoted.reset_index()
3             .melt(id_vars="employee", value_name="hire_date"))
```

- How do you make reverse look like df6 again?

Output

- ▶ https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html

Output

- ▶ https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html

```
1 df.to_csv(FNAME, sep=";")
2 df.to_html(FNAME, decimal="," , justify="center")
3 df.to_stata(FNAME, write_index=False)
```
